

Homework 2 PUBH 7440

Jake Wittman

2/17/2020

1a) From Gelman 3.2, the posterior mean given a non-informative uniform prior on $(\mu, \log\sigma)$ is

$$p(u|y_i) \sim t_{n-1}(\bar{y}, \frac{s^2}{n})$$

so $p(u_c|y_i) \sim t_{31}(1.013, \frac{0.025^2}{32})$ and


$$p(u_t|y_i) \sim t_{35}(1.173, \frac{0.20^2}{36})$$

1b)

```
mu.c <- 1.013 + (0.025/sqrt(32))*rt(1000,31)
mu.t <- 1.173 + (0.20/sqrt(36))*rt(1000,35)

poster_diff <- mu.t - mu.c
hist(poster_diff, breaks = seq(-0.1, 0.4, 0.02))
```

Histogram of poster_diff



The 95% credible interval is [0.0949057, 0.2279495]

2)

$$\begin{aligned}
p(\mu, \sigma^2 | y) &\propto p(y | \mu, \sigma^2) p(\mu, \sigma^2) \\
&\propto (\sigma^2)^{-n/2} \exp\left(-\frac{(n-1)s^2 + n(\bar{y} - \mu)^2}{2\sigma^2}\right) \sigma^{-1} (\sigma^2)^{-(\nu_0/2+1)} \exp\left(-\frac{\nu_0\sigma_0^2 + \kappa_0(\mu - \mu_0)^2}{2\sigma^2}\right) \\
&\propto \sigma^{-1} (\sigma^2)^{-((\nu_0+n)/2+1)} \exp\left(-\frac{\nu_0\sigma_0^2 + (n-1)s^2 + \frac{n\kappa_0(\bar{y} - \mu_0)^2}{n+\kappa_0} + (n+\kappa_0)\left(\mu - \frac{\mu_0\kappa_0 + n\bar{y}}{n+\kappa_0}\right)^2}{2\sigma^2}\right)
\end{aligned}$$

Because of the above,

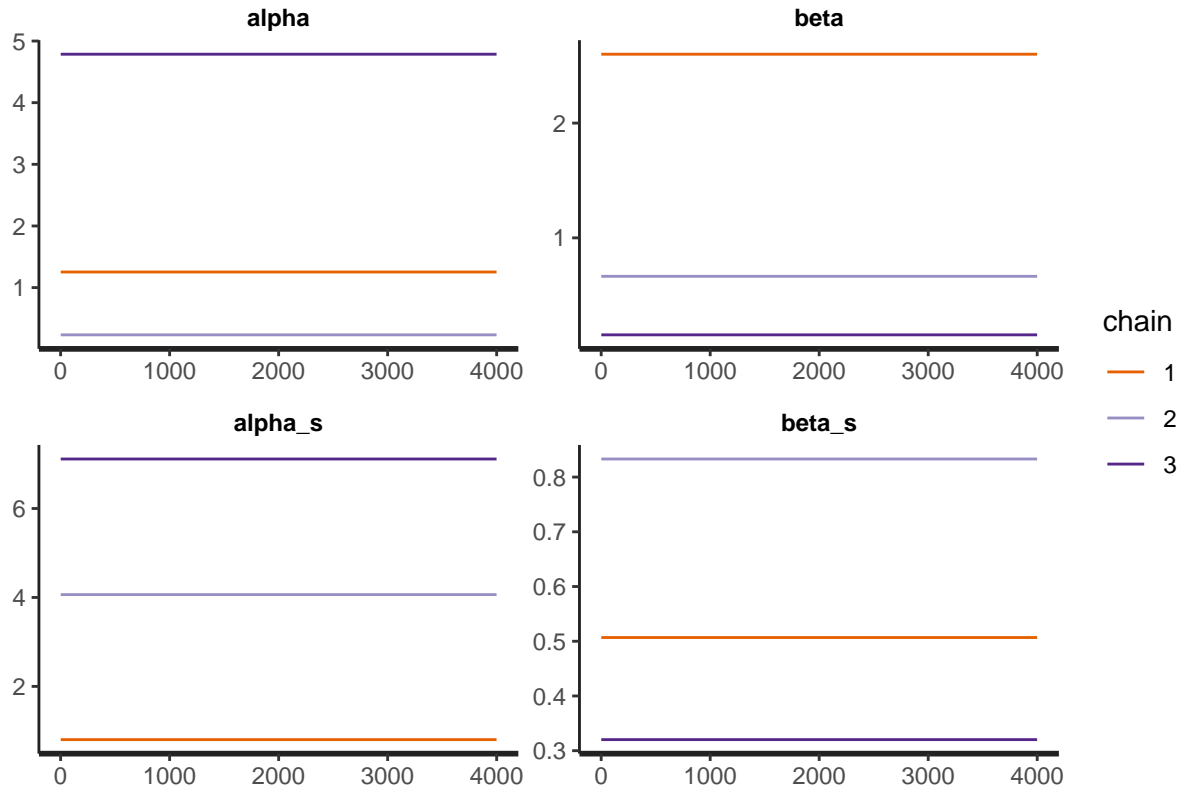
$$\mu, \sigma^2 | y \sim \text{N-Inv-}\chi^2\left(\frac{\mu_0\kappa_0 + n\bar{y}}{n+\kappa_0}, \frac{\sigma_n^2}{n+\kappa_0}; n + \nu_0, \sigma_n^2\right)$$


with

$$\sigma_n^2 = \frac{\nu_0\sigma_0^2 + (n-1)s^2 + \frac{n\kappa_0(\bar{y} - \mu_0)^2}{n+\kappa_0}}{n+\nu_0}$$


3a) Stan code used to fit the models is in the code appendix at the end of this assignment

Plots of the chains






Posterior plots for the difference in the population means and the new predicted y value.




Posterior plots with 95% credible intervals.


Posterior distribution for difference in population level means
with median and 95% credible intervals




Predictive distribution for the number of bikes on a new bike route intersection
with median and 95% credible intervals




4)






The final estimate for μ from importance sampling is 0.0035294. But I don't think I did this part right...



N(0,4.2) by rejection sampling



Code Appendix

```
knitr::opts_chunk$set(echo = FALSE)
mu.c <- 1.013 + (0.025/sqrt(32))*rt(1000,31)
mu.t <- 1.173 + (0.20/sqrt(36))*rt(1000,35)

poster_diff <- mu.t - mu.c
hist(poster_diff, breaks = seq(-0.1, 0.4, 0.02))
data {
  // Define data in this block

  //data for bike route intersections
  int<lower=0> n_s; // sample size for bike route intersections
  int<lower=0> y_s[n_s]; // bike route data

  //data for non-bike route intersections
  int<lower=0> n; // sample size for streets intersections
  int<lower=0> y[n]; // streets data
}

parameters {
  // Define parameters in this block
```

```

// parameters for bike route intersection
real<lower = 0> lambda_s[n_s]; //lambda for streets
real<lower = 0> alpha_s; // alpha for streets
real<lower = 0> beta_s; // beta for streets

// parameters for non-bike route intersection
real<lower = 0> lambda[n]; //lambda for streets
real<lower = 0> alpha; // alpha for streets
real<lower = 0> beta; // beta for streets
}

// The model
model {
  // Define model in this block

  // Model for bike route intersections
  target += gamma_lpdf(alpha_s | 0.01, 0.01); // gamma hyperprior on alpha
  target += gamma_lpdf(beta_s | 0.01, 0.01); // gamma hyperprior on beta
  target += gamma_lpdf(lambda_s | alpha_s, beta_s); // gamma prior on lambda
  target += poisson_lpmf(y_s | lambda_s); // distribution of response given lambda

  // Model for non-bike route intersections
  target += gamma_lpdf(alpha | 0.01, 0.01); // gamma hyperprior on alpha
  target += gamma_lpdf(beta | 0.01, 0.01); // gamma hyperprior on beta
  target += gamma_lpdf(lambda | alpha, beta); // gamma prior on lambda
  target += poisson_lpmf(y | lambda); // distribution of response given lambda
}

generated quantities{
  // Define other generated quantities in this block

  // Difference in lambdas
  real posterior_mean_difference;
  real pred_lambda;
  real pred_y;

  posterior_mean_difference = gamma_rng(alpha_s, beta_s) - gamma_rng(alpha, beta);

  pred_lambda = gamma_rng(alpha_s, beta_s);
  pred_y = poisson_rng(pred_lambda);
}

# Load libraries -----
pacman::p_load(tidyverse,
               rstan,
               brms,
               bayesplot)
# data -----
intersections <- c(1:18)

```

```

streets <- c(rep(1, 10), rep(0, 8))
bikes <- c(16, 9, 10, 13, 19, 20, 18, 17, 35, 55, 12, 1, 2, 4, 9, 7, 9, 8)

bike <- data.frame(
  intersections = intersections,
  streets = streets,
  bikes = bikes
)

bike_street <- bike %>% dplyr::filter(streets == 1)
bike_no_streets <- bike %>% dplyr::filter(streets == 0)

# Stan model -----
bike_model <- rstan::stan_model(file = here::here("homework/hw2.stan"))
# 3a -----
set.seed(1)
out_model1_streets <- sampling(
  object = bike_model,
  data = list(y_s = bike_street$bikes,
              n_s = nrow(bike_street),
              y = bike_no_streets$bikes,
              n = nrow(bike_no_streets)),
  warmup = 0,
  iter = 4000,
  chains = 3,
  cores = 4,
  control = list(adapt_delta = 0.9),
  show_messages = FALSE
)

traceplot(out_model1_streets, par = c("alpha", "beta", "alpha_s", "beta_s"))

# 3b -----
set.seed(1)
out_model1_streets <- sampling(
  object = bike_model,
  data = list(y_s = bike_street$bikes,
              n_s = nrow(bike_street),
              y = bike_no_streets$bikes,
              n = nrow(bike_no_streets)),
  warmup = 4000,
  iter = 6000,
  chains = 3,
  cores = 4,
  control = list(adapt_delta = 0.9),
  show_messages = FALSE
)

traceplot(out_model1_streets, par = c("alpha", "beta", "alpha_s", "beta_s"))
stan_hist(out_model1_streets, pars = c("posterior_mean_difference", "pred_y"))

```

```

mcmc_areas(out_model1_streets,
  pars = c("posterior_mean_difference"),
  prob = 0.95) +
labs(title = "Posterior distribution for difference in population level means",
  subtitle = "with median and 95% credible intervals",
  x = "Difference in population level means",
  y = "Posterior density") +
scale_x_continuous(limits = c(-50, 100),
  labels = seq(-50, 100, 25),
  breaks = seq(-50, 100, 25))

mcmc_areas(out_model1_streets,
  pars = c("diff_ys"),
  prob = 0.95) +
labs(title = "Predictive distribution for the number of bikes on a new bike route intersection",
  subtitle = "with median and 95% credible intervals",
  x = "Predicted number of bikes at a new bike route intersection",
  y = "Posterior density") +
scale_x_continuous(limits = c(0, 100),
  labels = seq(0, 100, 25),
  breaks = seq(0, 100, 25))

# Question 4 -----

# Importance sampling -----
posterior <- function(x) {

  post <- (1/3)*(1/sqrt(2*pi))*(exp((-1/2)*(x+5)^2) + exp((-1/2)*x^2) + exp((-1/2)*(x-5)^2))
  return(post)
}

# prep2: define the function to calculate log(weights)=log(p*(theta/x))-log(g(theta))
weight <- function(mu_j, c_mean, c_sd){

  posterior(mu_j) - dnorm(mu_j, mean = c_mean, sd = c_sd)
}

IS <- function(X, c, N) {

  # first calculate all the parameters
  n <- length(X)
  c_mean <- c * 0
  c_sd <- c * 4.2
  mu_j <- X

  # step 1: generate N samples from the importance function g(theta)
  theta <- rnorm(N, c_mean, c_sd)

```

```

# step 2: use the defined functions to calculate log of weights
# We work in the log scale for the sake of computational stability
lw <- weight(mu_j, c_mean, c_sd)

# The following step is to avoid infinite values when taking exponential of lw
# Using the step we have weights:  $w = \exp(lw) = \exp(w.scaled) * \exp(\max(lw))$ 
# Since w are in both the numerator and the denominator,  $\exp(\max(lw))$  will be canceled out
lw.scaled <- lw - max(lw)

## Finally, calculate the importance sampling estimate
theta.hat <- mean(theta * lw.scaled) / mean(lw.scaled)

## Now, evaluate the importance function using the coefficient of variance
c.v <- sd(lw.scaled)/mean(lw.scaled)

return(c(theta.hat,c.v))
}

## consider a grid of c values
c.candidate <- seq(-10, 10, length=500)

## specify vectors to store theta.hat and c.v for each value of c.candidate
hats <- rep(NA, 500)
cvs_hats <- rep(NA, 500)

## calculate estimate theta.hat and coef of var for each value of c.candidate
set.seed(12345)
vals <- seq(-10, 10, length.out = 500)
for(i in 1:500) {
  OUT <- IS(X = vals, c = c.candidate[i], 200)
  hats[i] <- OUT[1]
  cvs_hats[i] <- OUT[2]
}

## Examine where coef. of variance is low

plot(c.candidate, cvs_hats, ylab="coefficient of variance for mu from importance sampling", xlab="c", col="red", type="l")

## target mean estimator

plot(c.candidate, hats, xlab="c", ylab="E(X) estimate", col=4, type="l")

```

```

set.seed(12345)
N <- 200

# Rejection sampling

val <- seq(-10,10,by=0.01)
p.star <- exp(posterior(val))
plot(val, p.star, type = 'l', ylim = c(0,1.5), xlim = c(-10,10))

M.cand <- c(13:18)
for(i in 1:length(M.cand)){
  lines(val,dnorm(val, mean = 0, sd = 5)*M.cand[i],col=i+1,lty=2)
}
legend('topright',legend=paste0('M=',as.character(M.cand)),col=2:5,lty=2,cex=0.5)

RS <- function(M,N) {

  val.samples <- NULL

  while(length(val.samples) < N) {

    val.star <- rnorm(200, mean = 0, sd = 5)
    u <- runif(1)
    ratio <- posterior(val.star) / (M * dnorm(val.star, mean = 0, sd = 5))

    if(u<ratio) val.samples <- c(val.samples,val.star)
  }

  return(val.samples)
}

M <- 15
val.samples <- suppressWarnings(RS(M, N))

hist(val.samples, freq = FALSE, main = "N(0,4.2) by rejection sampling",xlab = expression(val))
lines(val, dnorm(val, mean = 0, sd = 4.2), col = 2)

```