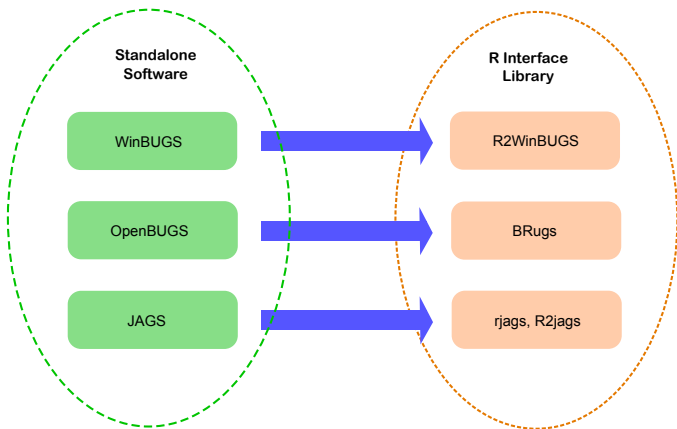


# R Interface for BUGS and Stan

Lin Zhang

Department of Biostatistics  
School of Public Health  
University of Minnesota

# Overview



**Pros:** user friendly GUI

**Cons:** cannot read data from file; tricky to do posterior inference

**Cons:** No interactive GUI

**Pros:** R interface facilitates reading data from file, posterior processing, plotting, etc.

# Relevant Softwares

To use the R implemented BUGS, we need to download and install the standalone softwares as well as the install the R packages in R:

- **R2WinBUGS:**
  - calls **WinBUGS** from within R
  - works in windows OS
- **BRugs:**
  - calls **OpenBUGS** from within R
  - information about OpenBUGS: <http://www.openbus.info/w>
  - runs in Windows, Mac, and Linux (not efficient in Linux)
- **rjags:**
  - calls **JAGS** from within R
  - available at <http://mcmc-jags.sourceforge.net>
  - runs stably in Windows, Mac, and Linux

# Toy Example

- Suppose the simple linear regression model

$$y_i = \beta_0 + \beta_1 x_i + \epsilon_i, \quad \epsilon_i \sim N(0, \sigma^2)$$

- Thus the **likelihood** is:

$$y_i \sim N(\beta_0 + \beta_1 x_i, \sigma^2)$$

- Assign vague **prior**:

$$\pi(\beta_0) = 1$$

$$\pi(\beta_1) = 1$$

$$\pi(\sigma) = \text{Unif}(0.01, 100)$$

## WinBUGS code

```
model{
  for(i in 1:N) {
    y[i] ~ dnorm(mu[i],y_prec)
    mu[i] <- beta[1]+beta[2]*x[i]
  }
  beta[1] ~ dnorm(beta_mean[1],beta_prec[1])
  beta[2] ~ dnorm(beta_mean[2],beta_prec[2])
  y_prec <- 1/(sigma*sigma)
  sigma ~ dunif(0.01,100)
}

# Data
list(x = c( 0.0, 0.41, 0.41, ...),
      y = c(1.80, 1.85, 1.87, ...),
      N = 27,
      beta_mean = rep(0,2), beta_prec = rep(0.01,2))

# Inits
list(beta=c(1,1),sigma=1)
list(beta=c(0,0),sigma=0.5)
```

## Preparation: write model to .txt file

- Before running BUGS in R, first we need to write and save the Bayesian model in a .txt file.
- Three ways to do it:
  - Use the BRugs/R2WinBUGS function '**writeModel**'
  - Use the base R command '**writeLines**'
  - Directly write in a text editor
- Code ...

# BRugs: model built step by step

- Recall the steps in WinBUGS:
  1. load model
  2. load data
  3. compile with specified number of chains
  4. load initials
  5. 'update' for burn-in iterations
  6. set parameters of which posterior samples will be collected
  7. 'update' again to collect posterior samples
  8. check results
- BRugs in R has the corresponding command for each step
- Code ...

## BRugs: meta function

- BRugs has a meta function '**BRugsFit**' to run model in **one step**

```
>BRugs.res <- BRugsFit(modelFile=modelfile,  
  data=list(x=x,y=y,N=N,beta_mean=beta_mean,  
  beta_prec=beta_prec), inits=initList,  
  numChains=2, parametersToSave = c("beta","sigma"),  
  nBurnin=1000, nIter=10000, nThin=1,  
  DIC=TRUE )
```

- Summary statistics and DIC values are returned

```
>BRugs.res$Stats  
>BRugs.res$DIC
```

- You can also use the same functions to check results, extract posterior samples

```
>samplesStats("*")  
>samplesHistory("*")
```

- Code ...



## CODA samples

- Set the argument 'coda=TRUE', the function 'BRugsFit' will return posterior samples in **coda format**.

```
>codaSamples <- BRugsFit(modelFile=modelfile,  
  data=list(x=x,y=y,N=N,beta_mean=beta_mean,  
  beta_prec=beta_prec), inits=initList,  
  numChains=2, parametersToSave = c("beta","sigma"),  
  nBurnin=1000, nIter=10000, nThin=1,  
  coda=TRUE )
```

- The resulting 'codaSamples' has these indices:

```
>codaSamples[[chainIdx]][iteratIdx, paramIdx]
```

Posterior samples of certain parameter can be extracted and saved for future use

- We can use the '**CODA**' R package for MCMC diagnostics.

# Convergence diagnostics with 'CODA' package

- Trace plots

```
>traceplot(codaSamples)  
>traceplot(codaSamples[, 'beta[1] ')
```

- Gelman-Rubin diagnostic

```
>gelman.diag(codaSamples)  
>gelman.plot(codaSamples)
```

- Autocorrelation plots

```
>autocorr.plot(codaSamples)
```

- Cross-correlation

```
>crosscorr(codaSamples)
```

- Effective sample size

```
>effectiveSize(codaSamples)
```

# rjags

- Three steps for Bayesian Analysis

1. Create, initialize, and adapt the model:

```
>jagsModel <- jags.model(file=modelfile,  
  data=list(x=x,y=y,N=N,beta_mean=beta_mean,  
  beta_prec=beta_prec), inits=initList,  
  n.chains = 2, n.adapt = 1000 )
```

2. Burnin iterations with 'update' function

```
>update(jagsModel, n.iter=1000)
```

3. Collect posterior samples in 'coda format'

```
>codaSamples <- coda.samples(model=jagsModel,variable.names  
  = c("beta","sigma"), n.iter = 10000, thin = 1)
```

- Note the 'rjags' package automatically load 'coda' package. Thus 'coda' functions can be directly used for convergence diagnostics for outputs from 'coda.samples' function.

- DIC calculation in rjags:

```
>dic.samples(model=jagsModel,n.iter=10000,thin=1,type='pD')
```

- Code ...

# R2WinBUGS

- Similar to the meta function in package 'BRugs'

- Main function

```
>bugs.res <- bugs(data=list(x=x,y=y,N=N,beta_mean=beta_mean,  
  beta_prec=beta_prec), inits=initList,  
  parameters.to.save = c("beta","sigma"),  
  model.file=modelfile,  
  n.chains=2, n.iter=10000, n.burnin=1000, n.thin=1,  
  DIC = TRUE, # codaPkg=TRUE,  
  bugs.directory='C:/Program Files/WinBUGS14/')
```

- Check result

```
>print(bugs.res)  
>plot(bugs.res)
```

- Return **coda-format** outputs if set the argument 'codaPkg=TRUE'.

- Code ...

# Stan: Efficient MCMC for Bayesian Analysis

- Stan is a C++ library for Bayesian modeling and inference.
- Like BUGS and JAGS, Stan generates posterior samples given a user-specified hierarchical model and data → **No need to derive the posteriors and code up the MCMC algorithm.**
- **Differently**, Stan uses **Hamiltonian Monte Carlo (HMC)** and **No-U-Turn Sampler (NUTS)** to produce high dimensional proposals that are **accepted with high probability** without having to spend time tuning.
- Many interfaces: RStan (R), PyStan (Python), MatlabStan (MATLAB), CmdStan (shell, command-line terminal), ...

# RStan: R interface to Stan

- The **rstan** package allows for fitting Stan models from R
- RStan has inbuilt diagnostics to analyse the MCMC outputs
- Typical workflow:
  1. Write a hierarchical model using Stan language and save in a file with `.stan` extension
  2. Run the Stan program by a single call to the **stan** function in R
  3. Diagnose non-convergence of the MCMC chains
  4. Conduct inferences based on the posterior samples

# Simple Example: Eight Schools

- Data:

School	$y$	$\sigma$
<i>A</i>	28	15
<i>B</i>	8	10
<i>C</i>	-3	16
<i>D</i>	7	11
<i>E</i>	-1	9
<i>F</i>	1	11
<i>G</i>	18	10
<i>H</i>	12	18

- Random effects model:

$$\begin{aligned}y_j &\sim N(\theta_j, \sigma_j^2) \quad j = 1, \dots, 8 \\ \theta_j &\sim N(\mu, \tau^2)\end{aligned}$$

- Stan and R codes ...

## Stan codes

```
data {  
  int<lower=0> J;           // number of schools  
  real y[J];               // estimated treatment effects  
  real<lower=0> sigma[J];  // s.e. of effect estimates  
}  
parameters {  
  real mu;  
  real<lower=0> tau;  
  vector[J] eta;  
}  
transformed parameters {  
  vector[J] theta;  
  theta = mu + tau * eta;  
}  
model {  
  eta ~ normal(0,1);  
  y ~ normal(theta,sigma);  
}  
//model {  
//  target += normal_lpdf(eta | 0, 1);  
//  target += normal_lpdf(y | theta, sigma);  
//}
```



## R codes

```
> schools_data <- list(
+   J = 8,
+   y = c(28, 8, -3, 7, -1, 1, 18, 12),
+   sigma = c(15, 10, 16, 11, 9, 11, 10, 18)
+ )

> library(rstan)

> fit1 <- stan(
+   file = "schools.stan", # Stan program
+   data = schools_data,   # named list of data
+   chains = 2,            # number of Markov chains
+   warmup = 1000,        # number of warmup iterations per chain
+   iter = 2000,          # total number of iterations per chain
+   cores = 2,            # number of cores (could use one per chain)
+   refresh = 0           # no progress shown
+ )
```

# Convergence diagnosis and posterior summary

- Check the convergence and summarize the posterior distribution

```
> traceplot(fit1, pars = c("mu", "tau"), inc_warmup = TRUE)
> print(fit1, pars=c("theta", "mu", "tau", "lp_"),
+       probs=c(.025,.5,.975))
```

- Check the between-parameter cross-correlation

```
> pairs(fit1, pars = c("mu", "tau", "lp_"), las = 1)
```

- Extract the posterior samples

```
> para_samples <- extract(fit1, pars=c("theta", "mu", "tau",
+                                     "lp_"), inc_warmup=FALSE)
```

- Some diagnostic tools for Hamilton Monte Carlo algorithm

```
> sampler_params <- get_sampler_params(fit1, inc_warmup = TRUE)
> summary(do.call(rbind, sampler_params), digits = 2)
```